

Performance and Power Management of Virtualized Platform

A Seminar Report
submitted in partial fulfillment of the
requirements for the degree of
Master of Technology
by

Under the guidance of
Prof Varsha Apte

Abhishek Pratap Singh
143050077



Computer Science
Indian Institute of Technology Bombay
Mumbai 460076 (India)

Abstract

This report is about what is virtualization and power management in virtualized platforms. There are many organisations and communities working in this field but xen and kvm are milestones in this field. xen started an open source project to host multiple guest OSes with an approach of paravirtualization. The interesting thing about kvm is it is integrated with Linux kernel so its development is happening with development of Linux kernel. Virtualization has performance issues so hardware support was the need of virtualization. Nowadays Intel and AMD provide hardware support. kvm provides virtualization using this hardware support. “which virtualization platform is best” is a hot topic, there is a need of regular comparison among all virtualization options to find out which one is best in a particular situation or requirement. This report describes about comparison of virtualization platforms also. Virtualization provides Consolidation and VM migration strategy to power management but this facility (virtual platform) comes with complication to manage power in virtualized platform this report describes how to deal with power management in virtualized platform. In last section focus on coordination among different but simultaneous power management plans, and some improvement in scheduler to make fair CPU capacity allocation during frequency scaling.

Contents

0.1	Introduction	2
0.2	Virtualization Technologies	3
0.2.1	Virtualization	3
0.2.2	xen	3
0.2.3	kvm	4
0.3	Comparison among xen, kvm and other Virtualization technologies	4
0.3.1	Description of experiment done to find performance related data. .	5
0.3.2	Performance analysis	5
0.4	Power Management in Virtualized Platforms	5
0.4.1	Difference between Power Management and Virtual Power Management	6
0.4.2	Mechanisms to manage virtual power	6
0.4.3	Power Consumption Reduction	8
0.5	Challenges related to virtual power management	8
0.5.1	Different Purpose power management	8
0.5.2	Conflict among simultaneous Power Management for different Purpose	9
0.5.3	How can we coordinate different power management solutions managing power simultaneously	9
0.5.4	Unfairness to some VM due to frequency scaling	9
0.5.5	Introduction of different VM Scheduler	10
0.5.6	How can we avoid unfairness to VMs by making changes to scheduler	11
0.6	Conclusion	11

0.1 Introduction

In Data centres there is need of different servers and different application, all of them can not run on same hardware without affecting each other, so there is need of deploying separate hardware for each such server and application this is very inefficient , costly and hard to manage. There is very nice solution for this problem is Virtualization. This is simple one of motivation for virtualization Technology. Virtualization provides a platform which create illusion to application or user that he is working with bare metal (hardware).

Virtualization should have some property to ensure feasibility and fairness to Virtual machines 1. Isolation if more than one VM running on a same physical machine then one VM should not affect adversely the performance of others. 2. support to various type of operating systems 3. overhead to provide virtualization should be least [1].

xen is a hypervisor which provides virtualization first released in 2003 by university of cambridge computer laboratory when xen was released hardware support for virtualization was not common so xen used paravirtualization approach to implement virtualization. Paravirtualization is a technique in which host operating system source code and guest operating system source code undergoes in modification to create virtual platform .

kvm provides virtualization but it runs only on Linux systems but integration with linux provides growth in development with linux kernel. kvm depends on hardware support made available by intel and AMD , and virtual machines run in guest mode of processor .[4]

In next section Comparison among xen, kvm and other Virtualization technologies. **Power Management in Virtualized Platforms :** In virtual platform access to hardware directly is not provided normally . If we tries to give direct access to hardware then it create some serious problems like violation of isolation property and there will be no difference in privilege level of host and guest and host.

Challenges Related to Virtual Power Management : This report focuses mainly on two Challenges 1.Inconsistency when different power management act simultaneously in un-coordinated way [7]

2.when more than one VM running on a single physical machine and all VM allocated some fixed credits, lets consider some VMs with less credits running overloaded but physical machine is underloaded globally then DVFS detects this global underload and reduce frequency. This reduction in frequency affect adversely overloaded VMs [3].

0.2 Virtualization Technologies

0.2.1 Virtualization

Virtualization is a way to provide a platform which creates illusion of hardware machine on which applications run without knowing it, but as time passed its meaning has changed because virtualization has a inherent problem of performance .

To increase speed of execution there are two simple ways to solve this either provide hardware support to virtualization or make high performance hardware for host OS. first method is better than second one because it reduces emulating overhead on hypervisor so its speed improves, but in second way due to high performance you are not seeing effects of overhead .

Virtualization provides independence on hardware , a server running on virtual machine can replicate and migrate according to need,

so maintenance of server farm has become very easy and automatic .

Live VM migration provides a way to use hardware efficiently in terms of power in technical term it is known as consolidation.

0.2.2 xen

xen is first independent open source virtualization technique which uses paravirtualization.

xen purposed paravirtualization technique to implement virtual platform for guest OSes . In paravirtualization we need to modify some source code of host OS and guest OS. To ensure isolation they multiplex Guest OSes instead of process[1]. xen[1] Paper deals with x86 architecture which not virtualization-friendly so if you want to implement virtualization then performance will be reduced and complexity will be increased . That's why they approached paravirtualization to make it faster. In this paper they didn't fully virtualized I/O devices , they used a abstraction of devices which is quite similar to hardware . There is a hypervisor (xen itself the hypervisor) which works at higher privilege level and provide interface between host and guest and validate communication between Guest OS and Hardware. There is special guest Dom 0 which controls all operation related to communication to network devices and i/o devices.

cons: 1. Xen need to modify in host OS and Guest OS some time it is not feasible ,in case

of commercial products which don't reveal source code.

2. In Disk Management if we are not using reorder barriers (i think reorder barrier rarely used) then we can turn of scheduling of request at Guest OS because Xen and disk scheduler also provide scheduling algorithm to reduce response time . So it is redundant work.

0.2.3 kvm

This open source virtualization technology released in 2007 and continuously developing with linux kernel. This is not independent virtualization solution it is part of linux kernel and implemented in kernel module , so its architecture to provide virtualization is depends on linux kernel .

kvm uses linux kernel as hypervisor so it skip one layer of interface(layer between host and hypervisor). kvm architecture is very simple functionality is implemented in kernel modules and user can manage VMs using ioctl system calls .

A VM is treated as a process in kvm and “*virtual CPU implemented using threads*”[5] so there is no need of separate scheduler for VMs .

kvm is a part of kernel so it can reuse code and functionality of kernel , and runs in kernel space unlike other hypervisors .

0.3 Comparison among xen, kvm and other Virtualization technologies

xen and kvm are different from each other in architecture and features . Hardware support was not available at the time when xen realeased , so to implement some special operation in guest os which requires privilege and to reduce overhead xen used paravirtualization but now hardware support is available so xen provides both configuration full virtaulization and paravirtualization . Difference in architecture: in xen guests perform all network and i/o operation with help of Dom 0, Dom 0 is a special guest to manage normal guests. but kvm is like extension of linux kernel so it uses linux kernel architecture as a hypervisor host and kvm uses same linux ”fair scheduler“ so managment of VM becomes very easy ,For example: user can pin VM to a particular core by using normal linux commands.

Features	KVM	XEN
architecture	Bare metal	Hosted
scheduler	(process scheduler) Linux fair schedulers versions	(VM scheduler) SEDF, Credit
Network Management	FIFO based scheduling	FIFO based scheduling
Memory Address Translation	Shadow page table, Hardware Assisted Pagetable	Direct Pagetable , shadow page table ,Hardware assisted pagetable

0.3.1 Description of experiment done to find performance related data.

[8]Experiments : all experiments done with high performance computing system
 Linpack: to test execution speed by solving linear system of equations , PingPong Bandwidth : measures bandwidth of communication among CPU cores ,Fast Fourier : measures the speed of execution of DFT, ping pong latency : measures latency of communication among CPU cores,SPEC OpenMP benchmark .

0.3.2 Performance analysis

High performance Computing :

kvm got highest rating point according to [8] and virtual box is very close to kvm , xen left behind in rating.

In ping pong bandwidth virtual box shows results better than native because virtual box doesn't pin its vcpu to a particular hardware cpu so there is a case when two or more vm running on same physical cpu that benefits in communication.

Normal Performance system :

[5] carried out some experiments CPU intensive , Memory intensive ,Memory Bomb and conclude both kvm and xen performed approximately equal .

0.4 Power Management in Virtualized Platforms

reference [6] provided conceptual idea about this topic.

Power management is directly related to hardware and to scale frequency there is need to communicate with hardware directly but direct access rights to VM violates crucial properties of virtualization like independence on hardware and isolation . So [6] provided nice solution for this use VPM states .

VPM states (soft states) are virtual version of p states provided by intel and these states are corresponding to performance

To implement local power management at VM level , VM continuously monitors utilization and according to that they requests desirable soft state and these request for change in soft state sometimes can be simply ignored by Dom0 to avoid control on hardware . all these request forwarded to Dom 0 , with the help of VM rules Dom0 takes appropriate action .

VPM rule are rules and constraints to define policies .

0.4.1 Difference between Power Management and Virtual Power Management

Datacentres use virtualization to make efficient use of hardware and for easy management of servers.

In this scenario Power management at host level must be coordinated with VMs. otherwise it will lose opportunities of global power management for eg. consolidation so it needs track details about load on VMs and load on physical machine .

0.4.2 Mechanisms to manage virtual power

Hardware Scaling

In this strategy hardware support required , nowadays intel hardware provides p states , p stands for performance these states provides the range of states in which lower states are high performance and consumes more power and higher states are corresponding to low performance and low power consumption .

These different states maintained by varying frequency of CPU. So using this hardware support power management system can manage power consumption by using appropriate p state according to current utilization of CPU.

Soft Scaling

When hardware support is limited for example cpu operate only on 2 or 3 different frequencies so to make significant reduction power consumption there is need of range of big range of p states .

In this situation we can exploit scheduling process in hypervisor to emulate more virtual p states, hypervisor scheduler can assign time slice according to soft state of that VM For

example if any VM requests to scale down frequency by 1/3 but CPU does not operate on frequency/3 here we can use scheduler to emulate this situation by assigning time slice which is one third of original time slice[6]

There is a big draw back in this approach that if a memory bound process gets less time slice then it affects adversely because it is not equivalent to frequency scaling operation ,if frequency scaling available then memory bound process can get CPU for short period of time and goes back to i/o operation but if time slice is reduced then it get stuck waiting for CPU meanwhile it could have done i/o instead of waiting for cpu,[6] doesnt provide effective solution for this condition.

Consolidation

This method is most effective to reduce power consumption and it operates globally . A central unit tracks all VMs load and tries to run VMs in a minimal set of hardware machine and turn down idle hardware machine . this strategy can be implemented easily by keeping records and tracking VMs load and their Hardware machine's load and search for a machine which can accommodate more VMs simultaneously if such Hardware machine found using live migration feature Vm can transfer from one machine to other machine .

Using C states

[2] provided basic idea of this topic.

When an system is idle it still consume a lot power and if goes for naive solution for this problem which is turn off idle systems then due to frequent changes in states of system idle to busy , busy to idle . so its not feasible to turn off system it will lead to drastic performance degradation because of large latency of turn on a system.

So we can opt second approach is a hardware support "C states", c states stands for core power states c0 is operating state other are non-executing states sleeping states varies in degree of sleeping.

Deeper states save more power but comes with drawback long latency to switch back in executing mode.

Actually these state implemented using turn of clocks and timers which are running in idle state of cpu .There is a Challenge with deeper c states . Deeper states turn off clocks and timer so there is problem of timer skew .

In this approach we can use tracking history of c states uses and prepare a c state plan based on the prediction .This approach is effective in periodic type behavior.

0.4.3 Power Consumption Reduction

Using all above mechanism reduction in power consumption can be done effectively all mechanism discussed above works on different different for example

Hardware Scaling and Soft Scaling acts on local level with DVFS each VM tries minimize power consumed .

Consolidation works at global level which tracks VMs and physical machines tries to accommodate max no. VMs on the minimal set of physical machine without affecting performance .

C states dealing with power consumption when cpu is idle .

so applying all these approaches we can reduce power consumption effectively.

0.5 Challenges related to virtual power management

0.5.1 Different Purpose power management

[7] provided the conceptual idea about this topic

There are different power management which serves different goals . If a power demand goes high and it remains high then there is risk of thermal fail-over which will destroy hardware this must not happen so there should be a plan to avoid thermal fail-over even in overloaded situation .

Average Power Consumption

This plan is to reduce average power consumption by using local and global power management.

Power Capping

This plan is to prevent thermal fail-over by restricting peak power consumption ,as observed running system on peak power for very short period of time does not lead to thermal fail-over , so a soft power management is fine .

Power planning on the basis of history

This plan works when behavior of load is periodic and predictable for such situation this plan tries to ensure that CPUs provide max resource s well as doesn't cross max peak power at the time of peak load , it can be done by a simple idea that before getting high load tries to finish previous work by giving some extra power .

0.5.2 Conflict among simultaneous Power Management for different Purpose

[7] provided the conceptual idea about this topic

When two or more different power management trying to achieve different goals there is great chance of Inconsistency. For example two power management plans working on same system one is for managing average power consumption and one is for managing max power and power budget .

both power management can interfere each other by overwriting p states of system.

There is one more scenario in which problem created due to lack of coordination between Consolidation plan running at global level and local power budget .Consolidation plan may accommodate VMs which violates power

0.5.3 How can we coordinate different power management solutions managing power simultaneously**0.5.4 Unfairness to some VM due to frequency scaling**

Assumption: There is no priority for any VM.

Lets take a Case of Service Provider and Consumer.

Consumer wants some computing resources (cpu) to run his VM but he wants to use a part of computation capacity so there is Credit allocation scheme which defines how much of computing resources will be available to consumer VM for example 40 credits means 40% of computing resources available to VM provided that cpu frequency is max .

Consider provider wants to reduce power consumption so there is power management which is using DVFS (Dynamic voltage and frequency scalling)

DVFS monitors utilization of cpu in closed feedback loop if utilization becomes lower then it scale down frequency to reduce power consumption

whenever this frequency scaled down there chances that some overloaded VM promised

X credit computing resources getting less amount of computing resources than promised because now frequency is scaled down. This creates unfairness toward VMs.

0.5.5 Introduction of different VM Scheduler

Fixed credit scheduler

In this type of VM scheduler if VM is promised X credit then VM will get at most X credit of computing resources even (Max credit-X) is unused and idle, and VM will get time slice according to credit X always.

Variable credit scheduler

In this type of VM scheduler if VM is promised X credit then VM will get X credit of computing resources provided that VM needs it if VM doesn't need, it is underloaded then free computing resources will be distributed among other needy VM. For example: Three VMs are running each allocated 33 credits. and Max credit is 99.

Two VMs are underloaded and one is overloaded then free time slices available from two underloaded VM to overloaded, **free** time slice means that two underloaded VM must get computing resources according to their need they must not suffer in terms of performance.

If all VM are at full utilization then all three VM will get promised (time slice) computing resources. No extra resources because computing resources are not free.

Did you notice that this variable credit scheduler solves "unfairness to VM" problem but there is a drawback in this solution that this is not desirable for service provider because unless all VM are underloaded or idle frequency scale down action can not be done. This is not desirable for provider because there are very less chance to scale down frequency which results in increased cost of power used, because power cost can not be minimized effectively in this case.

Ideal situation is when all VM gets computing resources as they promised always and no chance to scale down frequency should be prevented due to providing extra free time slice.

DVFS aware credit scheduler

[3] provided conceptual idea about this topic.

This scheduler considers both provider and consumer, this scheduler recomputes credit of

a VM on each frequency scale operation. For example if one VM allocated 30 credits and max frequency is F , if frequency scaled down to $F/2$ then VM getting actually 15% of computing resources so to make it equivalent to 30 credits at maximum frequency change credit 30 - 60 . Now 60 credit with $F/2$ frequency is equivalent to 30 credits.

0.5.6 How can we avoid unfairness to VMs by making changes to scheduler

[3] provided conceptual idea about this topic.

As DVFS aware credit scheduler described above it is different from other schedulers because it focus on Actual credit . Let say $RC = (\text{credit} \times \text{freq1})$ is real credit if frequency decreases multiplicatively by factor d .

Now new frequency is $\text{freq1}/d$ to keep RC constant credit should be multiplied by d , $RC = d \times \text{credit} \times \text{freq1}/d$ d cancel out and RC becomes again $\text{credit} \times \text{freq1}$.

0.6 Conclusion

As internet grows there is need of big cloud system and server farms , to manage easily and efficiently there is need of virtualization . Power consumption is directly related to environment so we are responsible to use it efficiently for this a good power management required. Combination of virtualization and power management is more effective .

References

- [1] Boris Fraser Keir Hand Steven Harris Tim Ho Alex Neugebauer Rolf Pratt Ian Warfield Andrew Barham, Paul Dragovic. Xen and the art of virtualization. pages 164–177. ACM SIGOPS Operating Systems Review, 2003.
- [2] Jiajun Xu Guanqun Lu Ke Yu Kevin Tian Gang Wei, Jinsong Liu. The on-going evolution of power management in xen.
- [3] Daniel Hagimont. Dvfs aware cpu credit enforcement in a virtualized system. pages 123–142, 2013.
- [4] Uri Liguori Anthony Kamay Yaniv Laor Dor Kivity, Avi Lublin. kvm: the linux virtual machine monitor. volume 1, pages 225–230. Proceedings of the Linux Symposium, 2007.
- [5] Sant Gadge Baba Amravati University Amravati India² Head Computer Centre Sant Gadge Baba Amravati University Amravati India Ms Jayshri Damodar Pagare 1, Dr. Nitin A Koli Research Scholar.
- [6] Karsten Nathuji, Ripal Schwan. Virtualpower: coordinated power management in virtualized enterprise systems. page 265, 2007.
- [7] Parthasarathy Talwar Vanish Wang Zhikui Zhu Xiaoyun Raghavendra, Ramya Ranganathan. No power struggles : Coordinated multi-level power management for the data center. volume 36, pages 48–59, 2008.
- [8] Andrew J. Younge, Robert Henschel, James T. Brown, Gregor von Laszewski, Judy Qiu, and Geoffrey C. Fox. Analysis of Virtualization Technologies for High Performance Computing Environments. *2011 IEEE 4th International Conference on Cloud Computing*, pages 9–16, 2011.